

Boltzmann Machines

Geoffrey E. Hinton

March 25, 2007

A Boltzmann Machine is a network of symmetrically connected, neuron-like units that make stochastic decisions about whether to be on or off. Boltzmann machines have a simple learning algorithm that allows them to discover interesting features in datasets composed of binary vectors. The learning algorithm is very slow in networks with many layers of feature detectors, but it can be made much faster by learning one layer of feature detectors at a time.

Boltzmann machines are used to solve two quite different computational problems. For a *search* problem, the weights on the connections are fixed and are used to represent the cost function of an optimization problem. The stochastic dynamics of a Boltzmann machine then allow it to sample binary state vectors that represent good solutions to the optimization problem.

For a *learning* problem, the Boltzmann machine is shown a set of binary data vectors and it must find weights on the connections so that the data vectors are good solutions to the optimization problem defined by those weights. To solve a learning problem, Boltzmann machines make many small updates to their weights, and each update requires them to solve many different search problems.

The stochastic dynamics of a Boltzmann machine

When unit i is given the opportunity to update its binary state, it first computes its total input, z_i , which is the sum of its own bias, b_i , and the weights on connections coming from other active units:

$$z_i = b_i + \sum_j s_j w_{ij} \tag{1}$$

where w_{ij} is the weight on the connection between i and j , and s_j is 1 if unit j is on and 0 otherwise. Unit i then turns on with a probability given by the logistic function:

$$\text{prob}(s_i = 1) = \frac{1}{1 + e^{-z_i}} \quad (2)$$

If the units are updated sequentially in any order that does not depend on their total inputs, the network will eventually reach a Boltzmann distribution (also called its equilibrium or stationary distribution) in which the probability of a state vector, \mathbf{v} , is determined solely by the “energy” of that state vector relative to the energies of all possible binary state vectors:

$$P(\mathbf{v}) = e^{-E(\mathbf{v})} / \sum_{\mathbf{u}} e^{-E(\mathbf{u})} \quad (3)$$

As in Hopfield nets, the energy of state vector \mathbf{v} is defined as

$$E(\mathbf{v}) = - \sum_i s_i^{\mathbf{v}} b_i - \sum_{i < j} s_i^{\mathbf{v}} s_j^{\mathbf{v}} w_{ij} \quad (4)$$

where $s_i^{\mathbf{v}}$ is the binary state assigned to unit i by state vector \mathbf{v} .

If the weights on the connections are chosen so that the energies of state vectors represent the badness of those state vectors as solutions to an optimization problem, then the stochastic dynamics of a Boltzmann machine can be viewed as a way of escaping from poor local optima while searching for good solutions. The total input to unit i , z_i , represents the difference in energy depending on whether that unit is off or on, and the fact that unit i occasionally turns on even if z_i is negative means that the energy can occasionally increase during the search, thus allowing the search to jump over energy barriers.

The search can be improved by using simulated annealing (Kirkpatrick et al., 1983). This scales down all of the weights and energies by a factor, T , which is analogous to the temperature of a physical system. By reducing T from a large initial value to a small final value, it is possible to benefit from the fast equilibration at high temperatures and still have a final equilibrium distribution that makes good solutions much more probable than bad ones. At a temperature of 0 the update rule becomes deterministic and a Boltzmann machine turns into a Hopfield net.

Learning in Boltzmann Machines

Given a training set of state vectors (the data), learning consists of finding weights and biases (the parameters) that make those state vectors good. More specifically, the aim is to find weights and biases that define a Boltzmann distribution in which the training vectors have high probability. By differentiating Eq. 3 and using the fact that $\partial E(\mathbf{v})/\partial w_{ij} = -s_i^{\mathbf{v}} s_j^{\mathbf{v}}$ it can be shown that

$$\sum_{\mathbf{v} \in \text{data}} \frac{\partial \log P(\mathbf{v})}{\partial w_{ij}} = \langle s_i s_j \rangle_{\text{data}} - \langle s_i s_j \rangle_{\text{model}} \quad (5)$$

where $\langle s_i s_j \rangle_{\text{data}}$ is the expected value of $s_i s_j$ in the data distribution and $\langle s_i s_j \rangle_{\text{model}}$ is the expected value when the Boltzmann machine is sampling state vectors from its equilibrium distribution at a temperature of 1. To perform gradient ascent in the log probability that the Boltzmann machine would generate the observed data when sampling from its equilibrium distribution, w_{ij} is incremented by a small learning rate times the RHS of Eq. 5. The learning rule for the bias, b_i , is the same as Eq. 5, but with s_j omitted.

If the observed data specifies a binary state for every unit in the Boltzmann machine, the learning problem is convex: There are no non-global optima in the parameter space. Learning becomes much more interesting if the Boltzmann machine consists of some “visible” units, whose states can be observed, and some “hidden” units whose states are not specified by the observed data. The hidden units act as latent variables (features) that allow the Boltzmann machine to model distributions over visible state vectors that cannot be modelled by direct pairwise interactions between the visible units. A surprising property of Boltzmann machines is that, even with hidden units, the learning rule remains unchanged. This makes it possible to learn binary features that capture higher-order structure in the data. With hidden units, the expectation $\langle s_i s_j \rangle_{\text{data}}$ is the average, over all data vectors, of the expected value of $s_i s_j$ when a data vector is clamped on the visible units and the hidden units are repeatedly updated until they reach equilibrium with the clamped data vector.

It is surprising that the learning rule is so simple because $\partial \log P(\mathbf{v})/\partial w_{ij}$ depends on all the other weights in the network. Fortunately, the difference in the two correlations in Eq. 5 tells w_{ij} everything it needs to know about the other weights. This makes it unnecessary to explicitly propagate error derivatives, as in the backpropagation algorithm.

Higher-order Boltzmann machines

The stochastic dynamics and the learning rule can accommodate more complicated energy functions (Sejnowski, 1986). For example, the quadratic energy function in Eq. 4 can be replaced by an energy function whose typical term is $s_i s_j s_k w_{ijk}$. The total input to unit i that is used in the update rule must then be replaced by $z_i = b_i + \sum_{j < k} s_j s_k w_{ijk}$. The only change in the learning rule is that $s_i s_j$ is replaced by $s_i s_j s_k$.

Conditional Boltzmann machines

Boltzmann machines model the distribution of the data vectors, but there is a simple extension for modelling conditional distributions (Ackley et al., 1985). The only difference between the visible and the hidden units is that, when sampling $\langle s_i s_j \rangle_{\text{data}}$, the visible units are clamped and the hidden units are not. If a subset of the visible units are also clamped when sampling $\langle s_i s_j \rangle_{\text{model}}$ this subset acts as “input” units and the remaining visible units act as “output” units. The same learning rule applies, but now it maximizes the log probabilities of the observed output vectors conditional on the input vectors.

Non-binary units (*not required reading*)

The binary stochastic units used in Boltzmann machines can be generalized to “softmax” units that have more than 2 discrete values, Gaussian units whose output is simply their total input plus Gaussian noise, binomial units, Poisson units, and any other type of unit that falls in the exponential family, which is characterized by the fact that the adjustable parameters have linear effects on the log probabilities (Welling et al., 2005). The general form of the gradient required for learning is simply the change in the sufficient statistics caused by clamping data on the visible units.

The speed of learning

Learning is typically very slow in Boltzmann machines with many hidden layers because large networks can take a long time to approach their equilibrium distribution, especially when the weights are large and the equilibrium distribution is highly multimodal, as it usually is when the visible units are unclamped. Even if samples from the equilibrium distribution can be obtained, the learning signal is very noisy because it is the difference of two sampled expectations. These difficulties can be overcome by restricting

the connectivity, simplifying the learning algorithm, and learning one hidden layer at a time.

Restricted Boltzmann machines

A restricted Boltzmann machine (Smolensky, 1986) consists of a layer of visible units and a layer of hidden units with no visible-visible or hidden-hidden connections. With these restrictions, the hidden units are conditionally independent given a visible vector, so unbiased samples from $\langle s_i s_j \rangle_{\text{data}}$ can be obtained in one parallel step. To sample from $\langle s_i s_j \rangle_{\text{model}}$ still requires multiple iterations that alternate between updating all the hidden units in parallel and updating all of the visible units in parallel. However, learning still works well if $\langle s_i s_j \rangle_{\text{model}}$ is replaced by $\langle s_i s_j \rangle_{\text{recon}}$ which is obtained as follows:

1. Starting with a data vector on the visible units, update all of the hidden units in parallel.
2. Update all of the visible units in parallel to get a “reconstruction”.
3. Update all of the hidden units again.

This efficient learning procedure does approximate gradient descent in a quantity called “contrastive divergence” and works well in practice (Hinton, 2002).

After learning one hidden layer, the activity vectors of the hidden units, when they are being driven by the real data, can be treated as “data” for training another restricted Boltzmann machine. This can be repeated to learn as many hidden layers as desired. After learning multiple hidden layers in this way, the whole network can be viewed as a single, multilayer generative model and each additional hidden layer improves a lower bound on the probability that the multilayer model would generate the training data (Hinton et al., 2006). Surprisingly, the resulting multilayer generative model is *not* a Boltzmann machine.

Learning one hidden layer at a time is a very effective way to learn deep neural networks with many hidden layers and millions of weights. Even though the learning is unsupervised, the highest level features are typically much more useful for classification than the raw data vectors. These deep

networks can be fine-tuned to be better at classification or dimensionality reduction using the backpropagation algorithm (Hinton and Salakhutdinov, 2006). Alternatively, they can be fine-tuned to be better generative models using a version of the “wake-sleep” algorithm (Hinton et al., 2006).

Relationships to other models (*not required reading*)

Boltzmann machines are a type of Markov random field, but most Markov random fields have simple, local interaction weights which are designed by hand rather than being learned. Boltzmann machines also resemble Ising models, but Ising models typically use random or hand-designed interaction weights.

The search procedure for Boltzmann machines is an early example of Gibbs sampling, a Markov chain Monte Carlo method which was invented independently (Geman and Geman, 1984) and was also inspired by simulated annealing.

Conditional random fields (Della Pietra et al., 1997) can be viewed as simplified versions of higher-order, conditional Boltzmann machines in which the hidden units have been eliminated. This makes the learning problem convex, but removes the ability to learn new features.

References

- Ackley, D., Hinton, G., and Sejnowski, T. (1985). A Learning Algorithm for Boltzmann Machines. *Cognitive Science*, 9(1):147–169.
- Della Pietra, S., Della Pietra, V., and Lafferty, J. (1997). Inducing features of random fields. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 19(4):380–393.
- Geman, S. and Geman, D. (1984). Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 6(6):721–741.
- Hinton, G. E. (2002). Training products of experts by minimizing contrastive divergence. *Neural Computation*, 14(8):1711–1800.

- Hinton, G. E., Osindero, S., and Teh, Y. W. (2006). A fast learning algorithm for deep belief nets. *Neural Computation*, 18:1527–1554.
- Hinton, G. E. and Salakhutdinov, R. R. (2006). Reducing the dimensionality of data with neural networks. *Science*, 313:504–507.
- Kirkpatrick, S., Gelatt Jr, C., and Vecchi, M. (1983). Optimization by Simulated Annealing. *Science*, 220(4598):671.
- Sejnowski, T. (1986). Higher-order Boltzmann machines. *AIP Conference Proceedings*, 151(1):398–403.
- Smolensky, P. (1986). Information processing in dynamical systems: Foundations of harmony theory. In Rumelhart, D. E. and McClelland, J. L., editors, *Parallel Distributed Processing: Volume 1: Foundations*, pages 194–281. MIT Press, Cambridge.
- Welling, M., Rosen-Zvi, M., and Hinton, G. E. (2005). Exponential family harmoniums with an application to information retrieval. In *Advances in Neural Information Processing Systems* **17**, pages 1481–1488. MIT Press, Cambridge, MA.